Tutorial on Image Clustering

by Joris Guérin

Which is the best CNN feature extractor?

Experiment description - datasets used

Natural object classification

VOC 2007

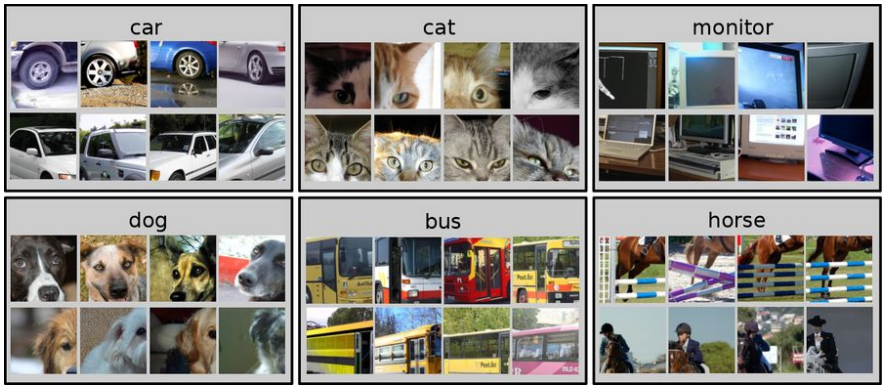

Coil 100

Experiment description - datasets used

Scene recognition

MIT 67

Architectural style (JTU, Shanghai)

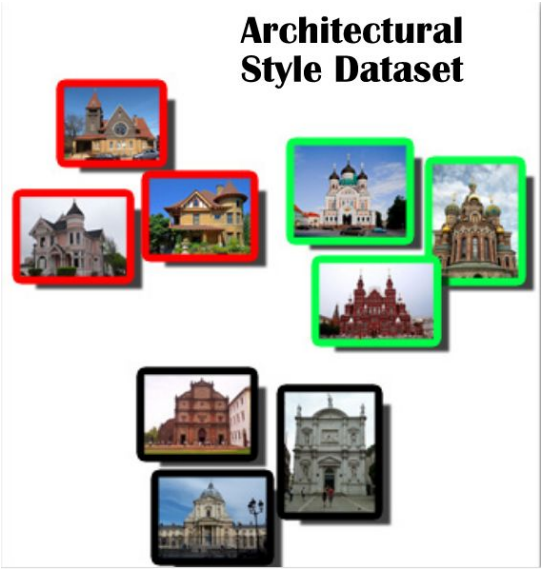

Experiment description - datasets used

Fine grained

Caltech UCSD Birds 200

Flower dataset (Oxford)

Experiment description - datasets used

Face recognition

UMist (Sheffield)

FEI (São Bernardo do Campo)

Experiment description - Experiments setup

Architectures used:
VGG16, VGG19, ResNet50, Inception, Xception

Layers used:

		VGG16	VGG19	Inception	Xception	Resnet50
L1	name	block5_pool	block5_pool	mixed7	add_12	activation_40
	shape	25,088	25,088	221,952	102,400	200,704
L2	name	fc1	fc1	mixed10	block14_sepconv2_act	activation_47
	shape	4,096	4,096	131,072	204,800	25,088
L3	name	fc2	fc2	avg_pool	avg_pool	avg_pool
	shape	4,096	4,096	2,048	2,048	2,048

Experiment description - Experiments setup

Clustering methods used:

KMeans and **Agglomerative clustering**

Metrics used:

NMI and **Purity**

Results summary

Layer choice:

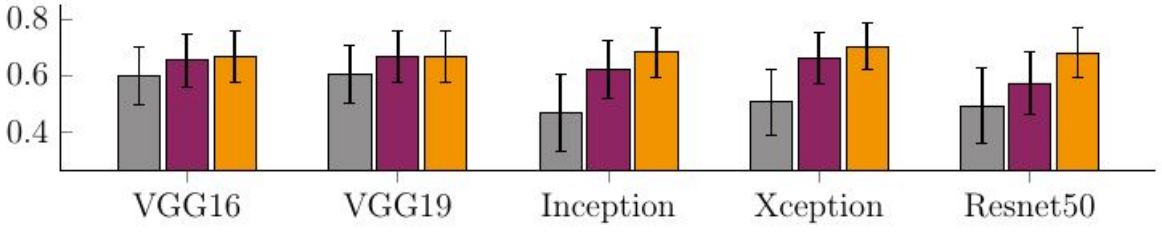

(a) Layer-architecture interaction
(mean and std across tasks and clustering algorithms).

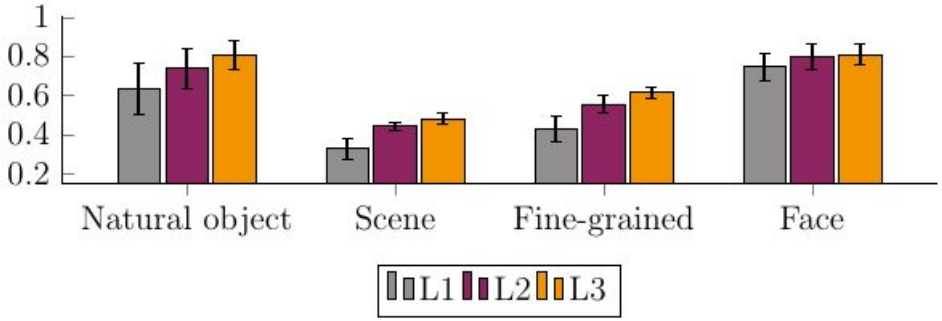

(b) Layer-task interaction
(mean and std across architectures, datasets and clustering algorithms).

Conclusion: **Use last layer**

Results summary

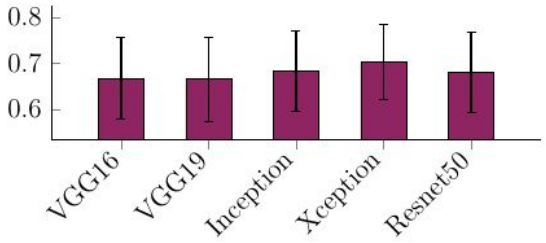

(a) mean and std across tasks and clustering algorithms

Architecture choice:

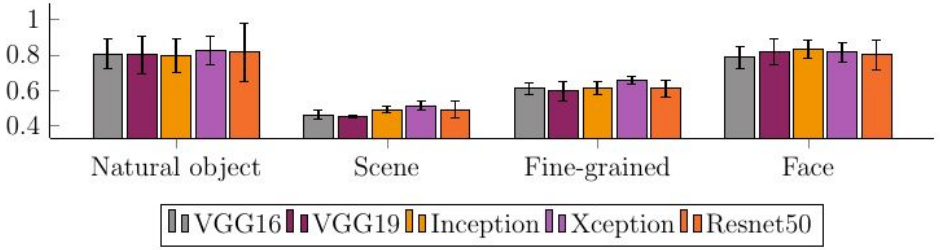

(b) Architecture-task interaction
(mean and std across datasets and clustering algorithms)

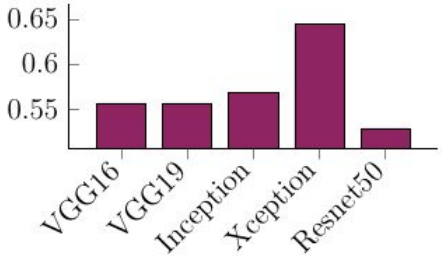

(a) Birds - Agglomerative clustering

(b) Flowers - Agglomerative clustering

Conclusion: **We don't know?!**

Possible strategies for new unsupervised dataset?

Selecting hyperparameters:

Supervised case:
Cross validation

Unsupervised case:
Follow the leader
(online learning)

Dimensionality reduction

Objectives and overview

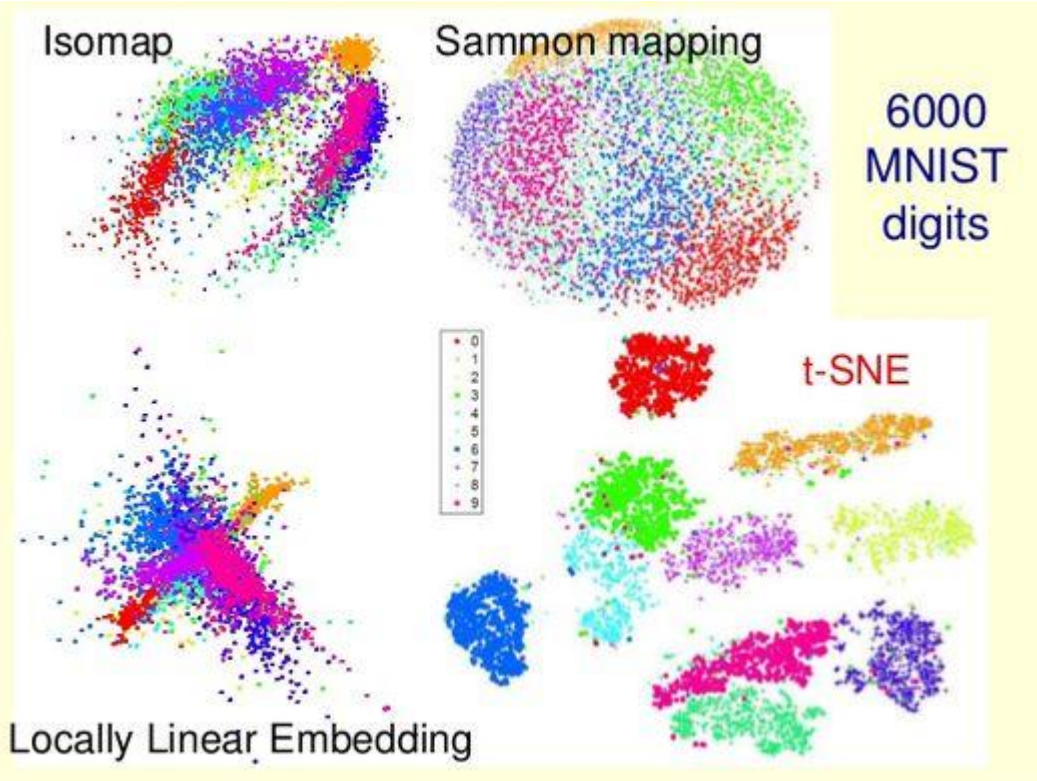

- Methods to visualize high dimensional data.
- Transform data into a 2D or 3D space

t-distributed Stochastic Neighbor Embedding (t-SNE)

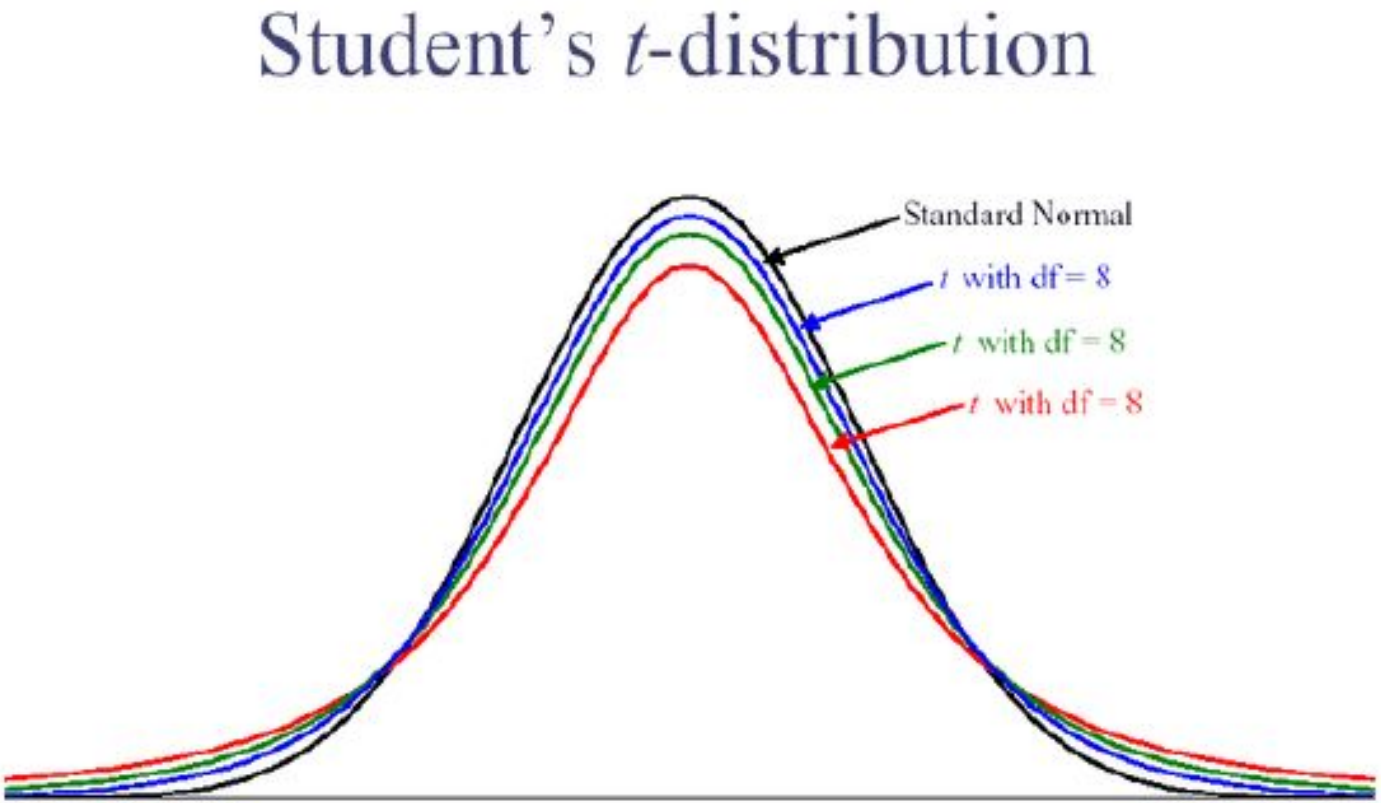

Ensemble of feature extractors

Why does it makes sense?

	NMI	PUR	FM	FM_{C_4}
InceptionResnet	**0.775**	**0.642**	**0.537**	0.442
VGG16	0.689	0.550	0.372	**0.653**
Densenet121	0.684	0.553	0.384	**0.700**

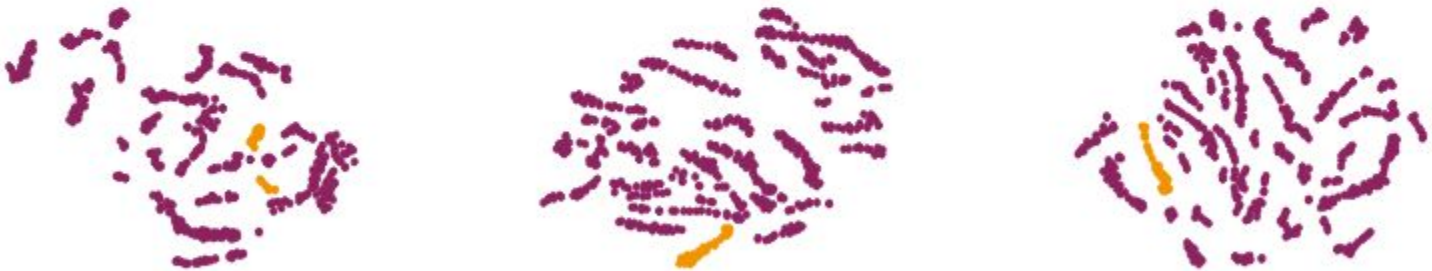

(a) InceptionResnet **(b)** VGG16 **(c)** Densenet121

Methodology

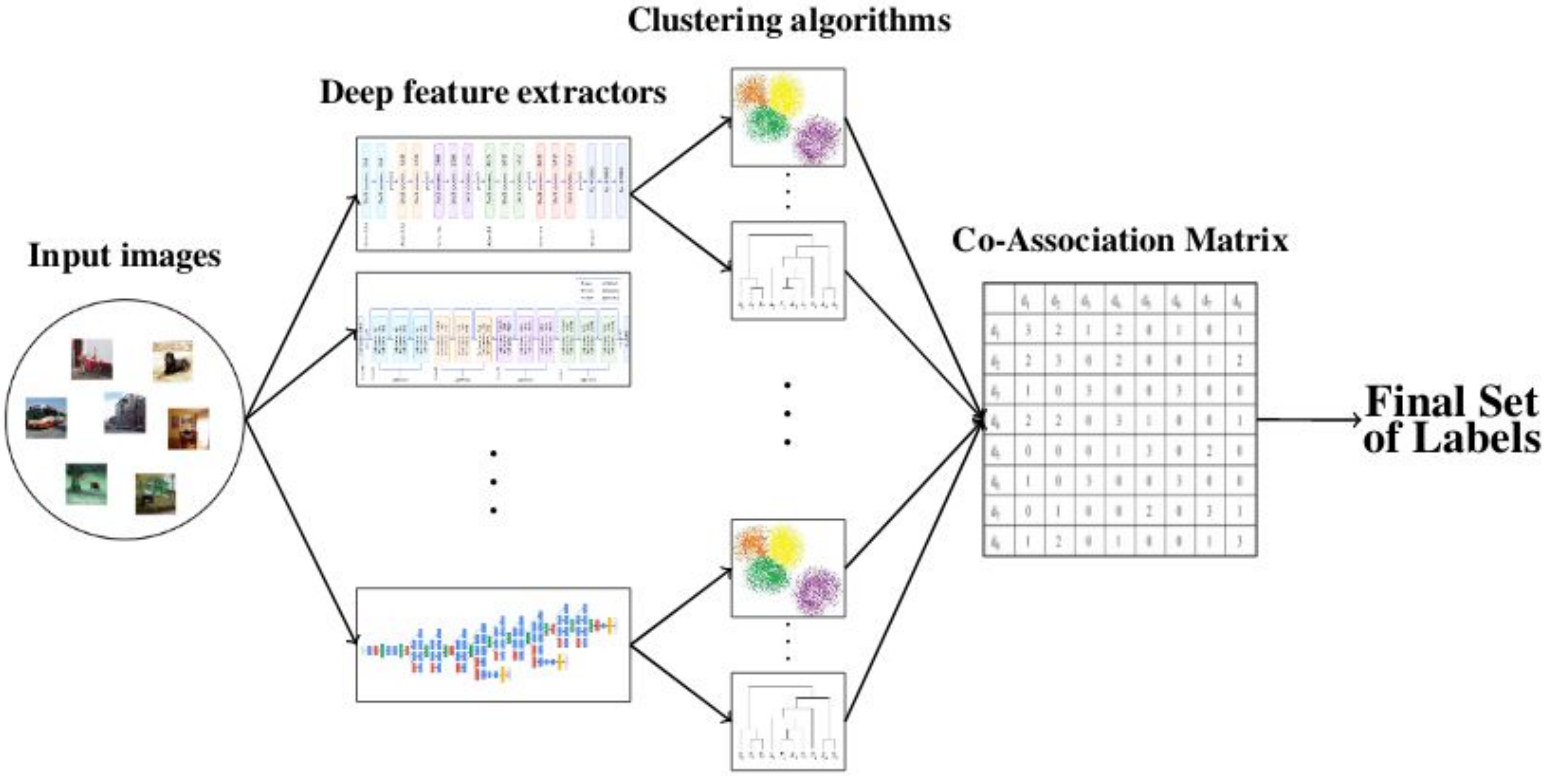

Results

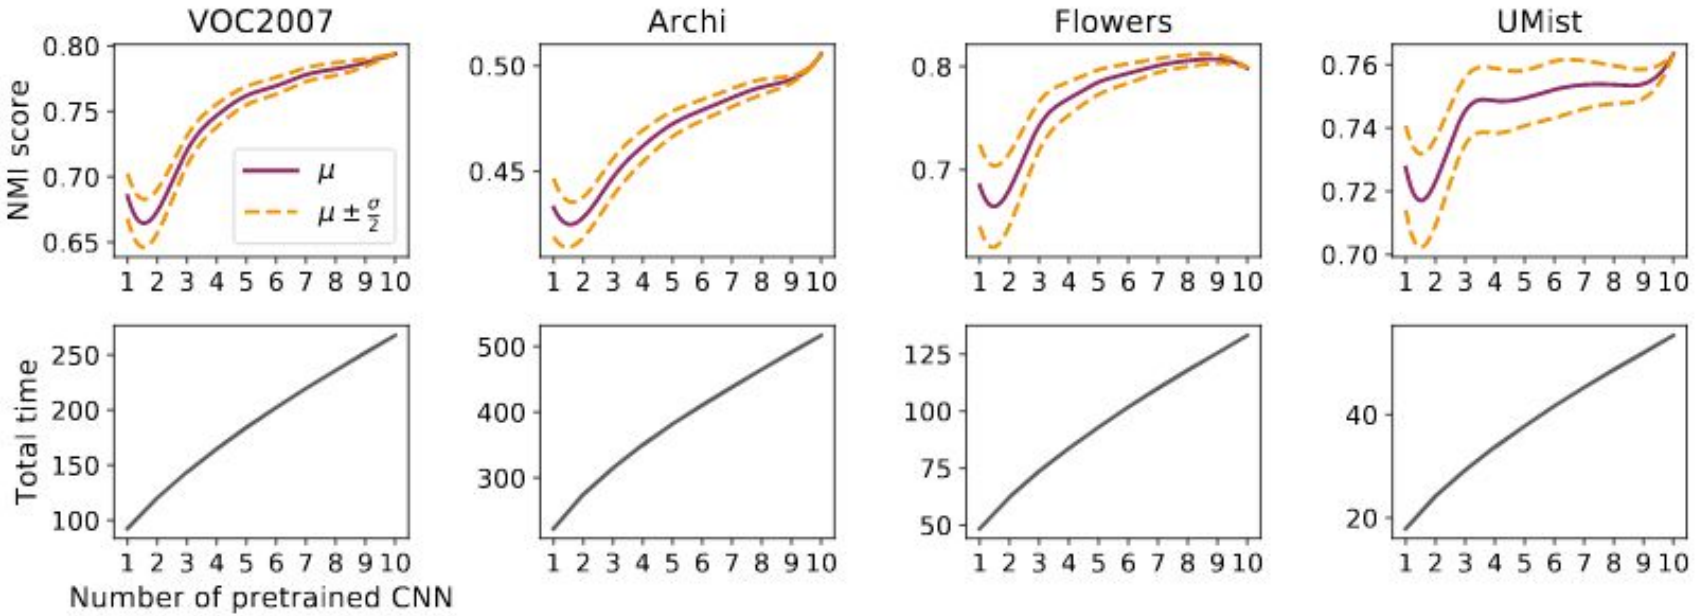
VOC2007
Archi
Flowers
UMist
NMI score
Total time
Number of pretrained CNN
μ
μ ± σ/2

Deep end to end clustering

Deep Embedded Clustering (DEC)

Unsupervised Deep Embedding for Clustering Analysis

Junyuan Xie JXIE@CS.WASHINGTON.EDU
University of Washington
Ross Girshick RBG@FB.COM
Facebook AI Research (FAIR)
Ali Farhadi ALI@CS.WASHINGTON.EDU
University of Washington

Objective: Learn jointly cluster assignment and new data representation

Autoencoder initialization

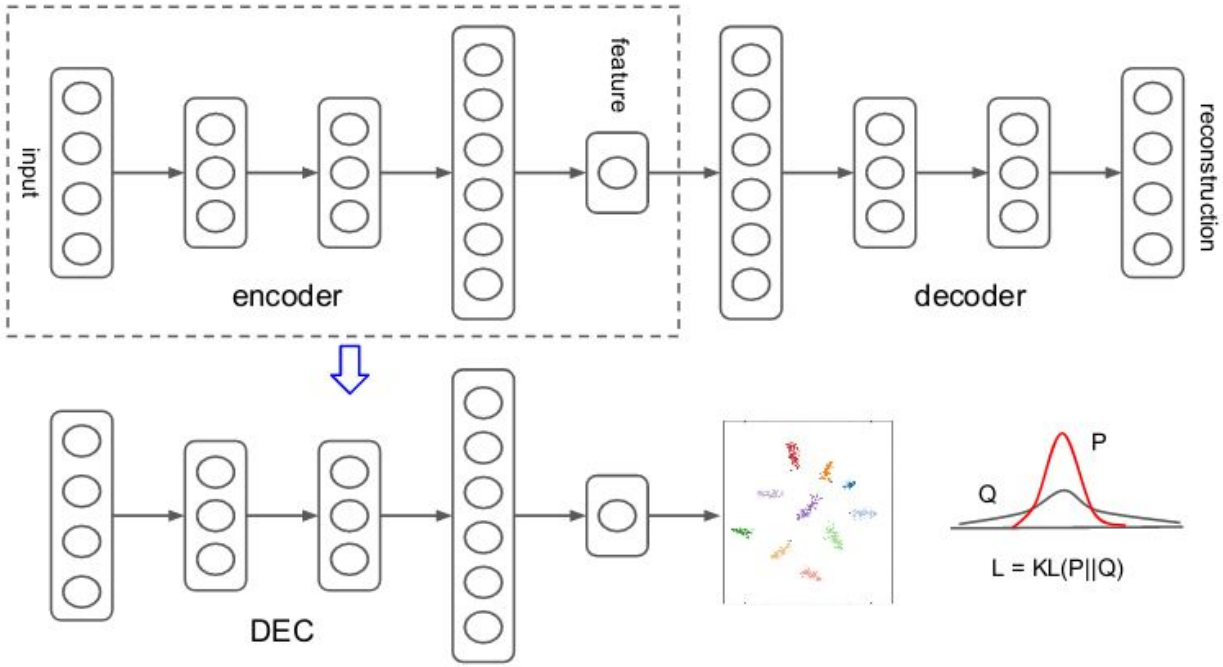

1- Denoising AE

2- End-to-end reconstruction AE

Joint optimization

- Soft assignment: Student t-distribtuion
(similarity between embedded point z_i and centroid μ_j)

$$q_{ij} = \frac{(1+\left\|z_i-\mu_j\right\|^2/\alpha)^{-(\frac{\alpha+1}{2})}}{\sum_{j'}(1+\left\|z_i-\mu_{j'}\right\|^2/\alpha)^{-(\frac{\alpha+1}{2})}}$$

- Target distribution: Strengthen high confidence prediction, normalize loss contribution of each cluster

$$p_{ij} = \frac{q_{ij}^2/f_j}{\sum_{j'} q_{ij'}^2/f_{j'}};\ f_j = \sum_j q_{ij}$$

- optimization: Gradient descent to update θ and μ_j to minimize

$$\mathcal{L} = KL(P||Q) = \sum_i \sum_j p_{ij} \log \frac{p_{ij}}{q_{ij}}$$

Results MNIST

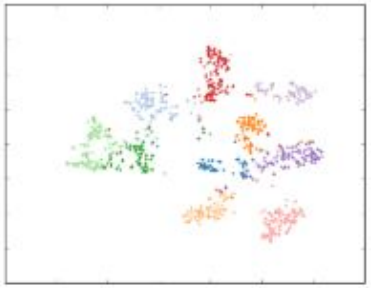

(a) Epoch 0

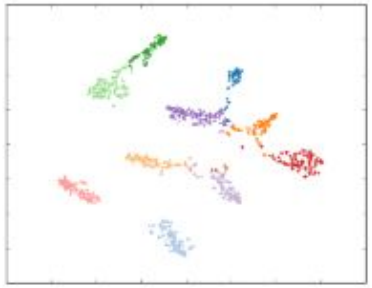

(b) Epoch 3

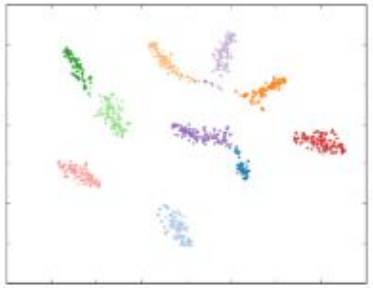

(c) Epoch 6

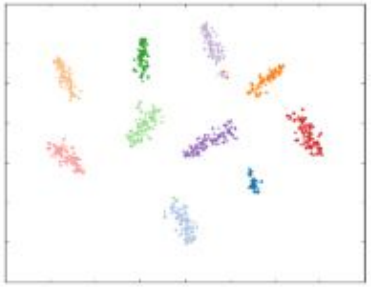

(d) Epoch 9

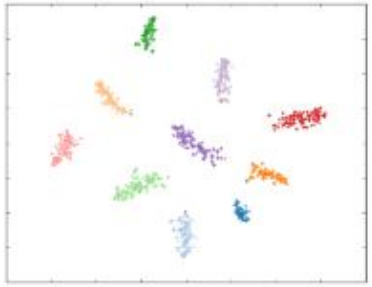

(e) Epoch 12

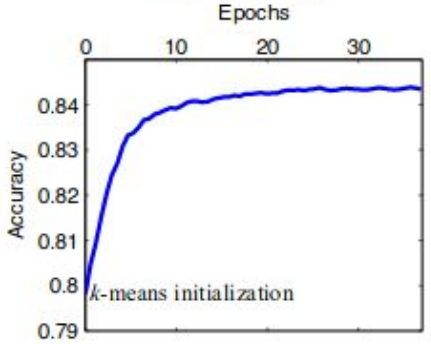

(f) Accuracy vs. epochs

JULE for networks gathering

Joint Unsupervised Learning of Deep Representations and Image Clusters

Jianwei Yang, Devi Parikh, Dhruv Batra
Virginia Tech
{jw2yang, parikh, dbatra}@vt.edu

Objective: Learn jointly cluster assignment and new data representation

Triplet loss

$$Loss = \sum_{i=1}^{N} \left[\|f_i^a - f_i^p\|_2^2 - \|f_i^a - f_i^n\|_2^2 + \alpha \right]_+$$

Schroff et al.

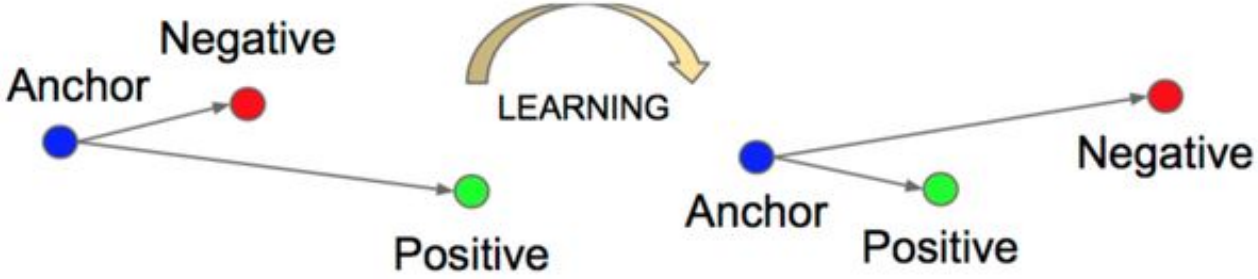

Schroff et al.

Deep Multi-view clustering

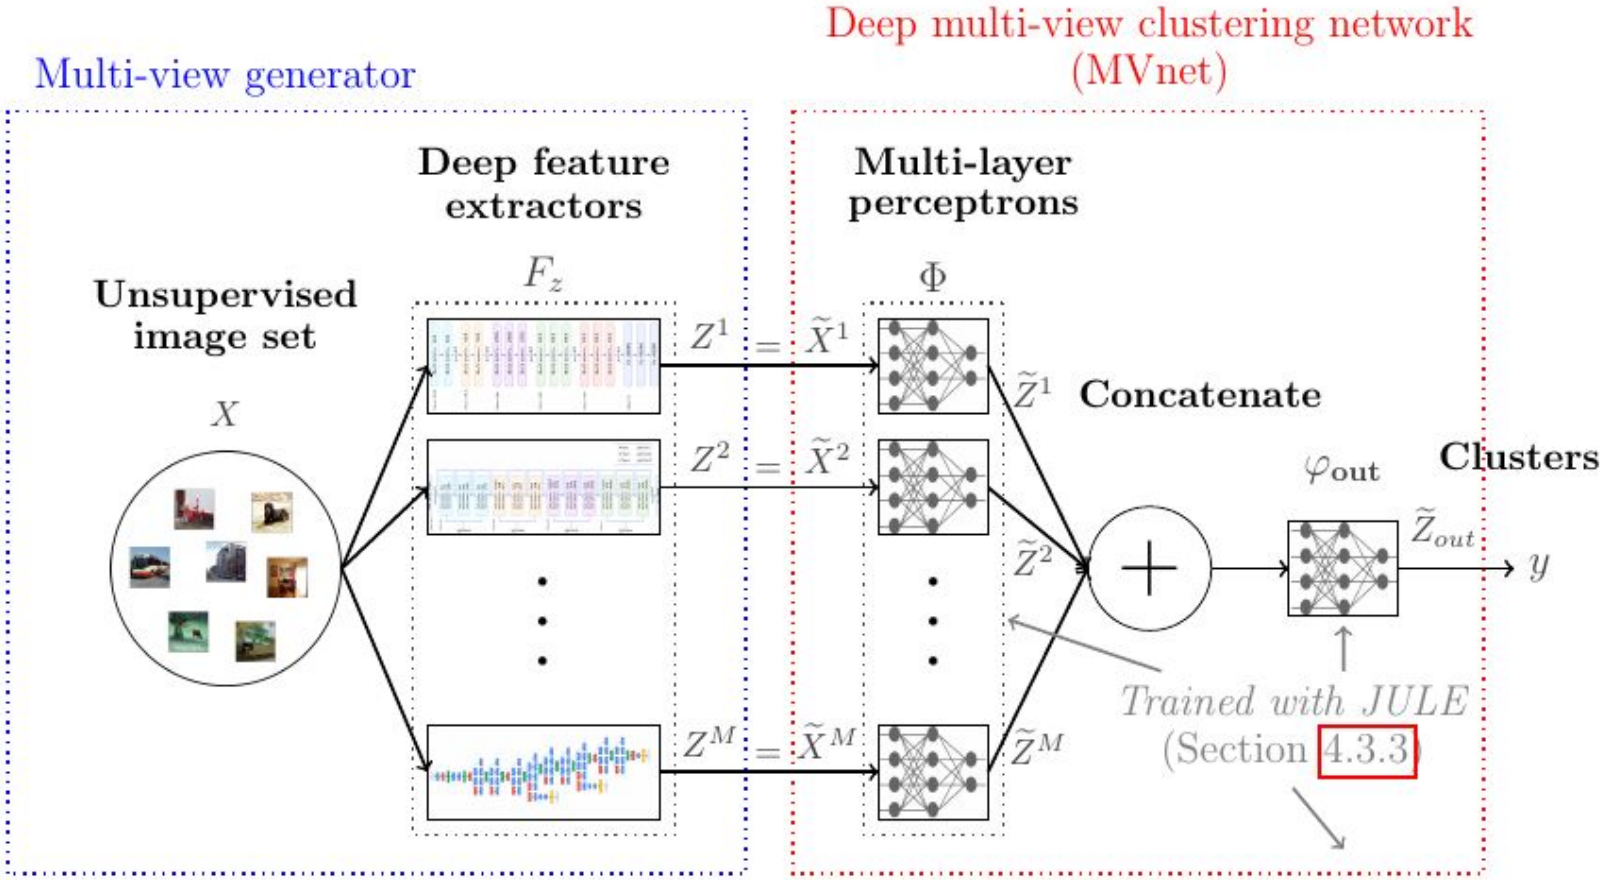

Results on UMist

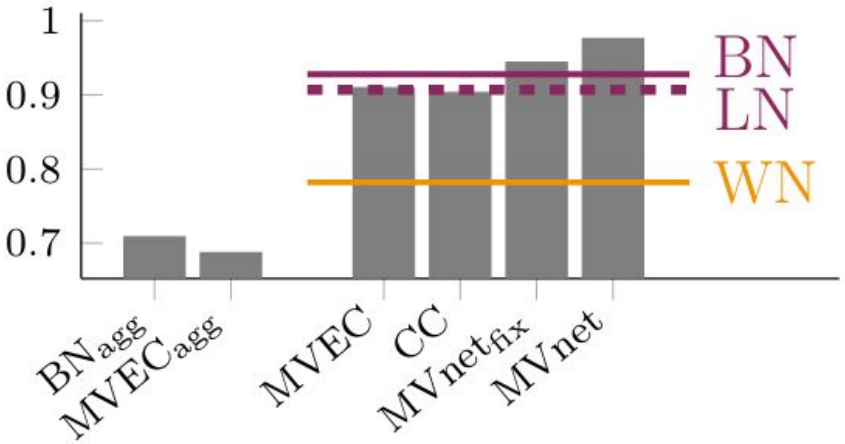

(g) UMist

Final scores:

Purity : 0.967
NMI : 0.984

(a) Densenet169 features

(b) Densenet169 + JULE

(c) Concat

(d) $\text{MVnet}_{\text{fix}}$

(e) MVnet

Conclusions

Multiple deep CNN feature extractors

+

Deep Clustering (JULE)

Initial scores:

Purity : 0.503
NMI : 0.663

Final scores:

Purity : 0.967
NMI : 0.984